

Impact of Agriculture on Global Warming

Tanmay Singh

Roll no. 2021569
tanmay21569@iiitd.ac.in

Ritwik Harit

Roll no. 2021557
ritwik21557@iiitd.ac.in

Shreyas Kabra

Roll no. 2021563
shreyas21563@iiitd.ac.in

Vasan Vohra

Roll no. 2021572
vasan21572@iiitd.ac.in

Abstract—The increasing global population demands growth in the agri-food industry. Understanding and addressing the environmental impact of the agri-food industry is crucial for mitigating climate change and developing sustainable practices within this sector.

The motivation for this project is the need to develop tools for predicting and managing the temperature change caused by CO₂ emissions. By harnessing the power of ML, the project aims to create a predictive model that will enable stakeholders and anyone connected to the agri-food industry to make informed decisions and reduce carbon emissions.

This paper's preprocessing techniques are handling Null values, Eliminating Duplicate Rows, Feature removal, One-Hot encoding, Standard Scaling, and Dimensionality Reduction. Model Validation is done using K-Fold cross-validation.

The Evaluation metrics used are Mean Absolute Error, Root Mean Square Error, and R² score. The models used are Random Forest Regressor, Linear Regressor, Gradient Boosting Regressor, Adaboost Regressor, XGB Regressor, Lasso Regression, Ridge Regressor, and SVR using various Kernels.

I. Introduction

In the recent Era, the agricultural sector has had a pivotal role in the significant contribution of greenhouse gasses and the potential for sustainability. The global architecture sector serves as a primary source of sustenance for billions of people on Earth but also leads to a high contribution in annual CO₂ emissions.

The climatic change due to the emissions depends on a lot of factors like Savanna fires, Forest fires, Crop Residues, Rice Cultivation, Drained organic soils (CO₂), Pesticides Manufacturing, Food Transport, Forestland, Net Forest conversion, Food Household Consumption, Food Retail, On-farm Electricity Use, Food Packaging, Agri-food Systems Waste Disposal, Food Processing, Fertilizers Manufacturing, IPPU (Industrial processes and product use), Manure applied to Soils, Manure left on Pasture, Manure Management, Fires in organic soils, Fires in humid tropical forests, On-farm energy use, Rural population, Urban population, Total Population - Male, Total Population - Female and total-emission.

The dataset and power of Machine Learning provide us with

tools to analyze the climatic and environmental effects of CO₂ emissions. This helps us to develop a road towards sustainable development. This provides invaluable insights to climate scientists, environmental researchers, and policymakers committed to mitigating the agricultural sector's environmental impact.

II. Literature Survey

Machine learning and deep learning have gained immense attention for their ability to unveil insights from data in diverse fields like image recognition, natural language processing, healthcare, agriculture, and more. Notably, they've played a pivotal role in climate science, especially in predicting global temperature changes.

In the paper titled "A Machine Learning-Based Model for Predicting Temperature Under the Effects of Climate Change" [2] by Mahmoud Y Shams et al. (2023), the authors explored global climatic patterns. Using the 'Climate Change' dataset, they employed various models such as Linear Regression, Random Forest, Decision Tree, K-Nearest Neighbor, Support Vector Machine, and Cat Boost Regressor. Their results indicated that the Cat Boost Regressor (CBR) Model outperformed others, achieving a 0.003 Mean Squared Error, 0.054 Root Mean Squared Error, 0.0036 Mean Absolute Error, and an R² score of 0.92.

Other significant studies include "Monthly prediction of air temperature in Australia and New Zealand with machine learning models" by S. Salcedo-Sanz et al. (2016) [3] and "Climate Change Analysis Using Machine Learning" by Himanshu Vishwakarma (2018) [4]. These studies explored temperature prediction using Support Vector Regressor (SVR), Multi-layer Perceptron (MLP), Linear Regression, SVR, Lasso Regression, and Elastic Net based on historical climate and greenhouse gases data.

III. Dataset Details

The Agri-food CO₂ emission dataset [1] available on Kaggle has been curated by combining and meticulously processing

multiple distinct datasets sourced from the Food and Agriculture Organization (FAO) and data provided by the Intergovernmental Panel on Climate Change (IPCC). As shown by the dataset, these emissions make a significant and noteworthy contribution to the annual global emissions.

The dataset contains 6965 rows and 31 columns, including 30 distinct features and 1 target column. The target column, labelled "Average Temperature °C ", indicates the yearly average temperature rise.

A. Key Features in the Dataset

- Area: Denotes the respective country.
- Year: Represents the specific year of data.
- Savanna Fires: Reflects emissions resulting from fires occurring in savanna ecosystems.
- Forest Fires: Represents emissions originating from fires within forested regions.
- Crop Residues: Signifies emissions from the combustion or decomposition of residual plant material post-harvesting.
- Rice Cultivation: Indicates emissions stemming from methane release during rice cultivation.
- Drained Organic Soils (CO₂): Quantifies emissions linked to carbon dioxide release during organic soil drainage.
- Pesticides Manufacturing: Quantifies emissions associated with pesticide production.
- Food Transport: Measures emissions resulting from the transportation of food products.
- Forestland: Specifies the land area covered by forests.
- Net Forest Conversion: Depicts changes in forest areas resulting from deforestation and afforestation.
- Food Household Consumption: Records emissions attributable to food consumption at the household level.
- Food Retail: Encompasses emissions generated by the operation of food retail establishments.
- On-farm Electricity Use: Represents electricity consumption on agricultural farms.
- Food Packaging: Quantifies emissions arising from producing and disposing of food packaging materials.
- Agri-food Systems Waste Disposal: Represents emissions emanating from waste disposal within the agri-food system.
- Food Processing: Represents emissions associated with the processing of food products.
- Fertilizers Manufacturing: Quantifies emissions linked to the production of fertilisers.
- IPPU (Industrial Processes and Product Use): Encompasses emissions originating from industrial processes and product use.
- Manure Applied to Soils: Reflects emissions from animal manure's application to agricultural soils.
- Manure Left on Pasture: Quantifies emissions arising from the presence of animal manure on pasture or grazing land.

- Manure Management: Indicates emissions related to the management and treatment of animal manure.
- Fires in Organic Soils: Captures emissions generated by fires occurring in organic soils.
- Fires in Humid Tropical Forests: Measures emissions resulting from fires in humid tropical forests.
- On-farm Energy Use: Represents energy consumption on agricultural farms.
- Rural Population: Signifies the number of individuals residing in rural areas.
- Urban Population: Represents the number of individuals residing in urban areas.
- Total Population - Male: Quantifies the total male population.
- Total Population - Female: Quantifies the total female population.
- Total Emission: Aggregates total greenhouse gas emissions from diverse sources.

B. Target Column

- Average Temperature °C: This column records the average annual temperature increase in degrees Celsius, serving as the dataset's primary target variable for analysis and prediction.

C. Preprocessing Techniques

- Handling Null Values: Rows containing Null values were removed from the dataset to ensure data integrity and prevent potential bias in the analysis.
- Eliminating Duplicate Rows: Duplicate rows were identified and removed from the dataset to avoid redundancy and ensure the accuracy of the analysis.
- Removed Outliers: Outliers were removed from the dataset because they have the potential to significantly impact the dataset by introducing variations and deviations from its typical distribution.
- Feature Removal: Features exhibiting a high correlation, represented by a correlation coefficient greater than or equal to 0.99, were identified and removed. This step helps in improving the model's interpretability and generalization.
- One-Hot Encoding for Categorical Features: The categorical feature "Area" was transformed using One-Hot encoding. This technique creates binary columns for each category, allowing machine-learning algorithms to work effectively with categorical data.
- Standard Scaling: Standard scaling was used to standardize the data. Standardization helps ensure that features with different units and scales do not disproportionately influence the modeling process and helps machine learning algorithms converge faster.
- Dimensionality Reduction: We utilized Principal Component Analysis (PCA) to reduce feature dimensionality from 183 to 160. We reduced the features to 160 because

the graph drawn for *Explained Variance vs Number of Components* (see Fig. 1) shows that 160 features are able to explain more than 95% of the variance in the data. This process improved efficiency, increased interpretability, and lowered computational complexity in our analysis.

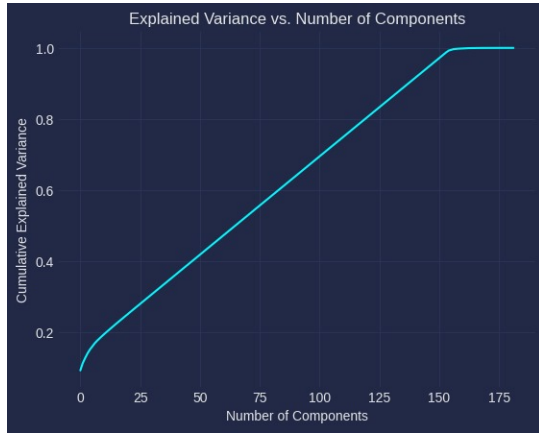


Fig. 1: *Explained Variance vs Number of Components*

IV. Methodology

A. Data Preprocessing

- Removed rows containing Null values.
- Eliminated duplicate rows.
- Created various graphs, including scatter plots, histograms, boxplots, pie charts, and t-SNE graphs, for data visualization and analysis.
- We removed features *IPPU* and *Total Population-Female* due to their high correlation (greater than or equal to 0.99) with *Food Household Consumption* and *Total Population-Male*, respectively, as identified through correlation heatmaps (see Fig. 2) and pair plots (see Fig. 3 and 4).
- Used 2 techniques for encoding the categorical features:
 - Label Encoding to transform the categorical feature into the numerical format.
 - Applied one-hot encoding to handle categorical features, increasing the number of features to 183. Subsequently, we attempted to reduce the feature set from 183 to 160 using PCA.

B. Model Validation

Utilized K-Fold cross-validation with K=5 to validate machine learning models. This approach helps assess the model's performance across different subsets of the data.

C. Evaluation Metrics

Used Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R2 score as evaluation metrics for both training

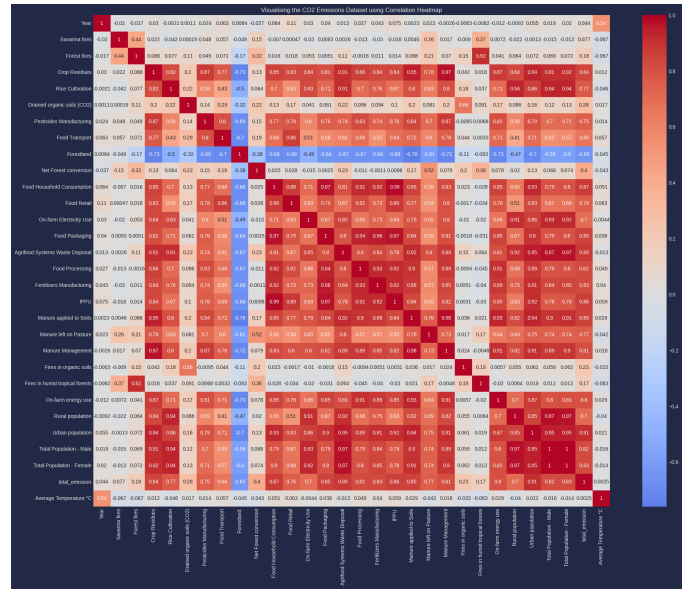


Fig. 2: *Correlation Heatmap*

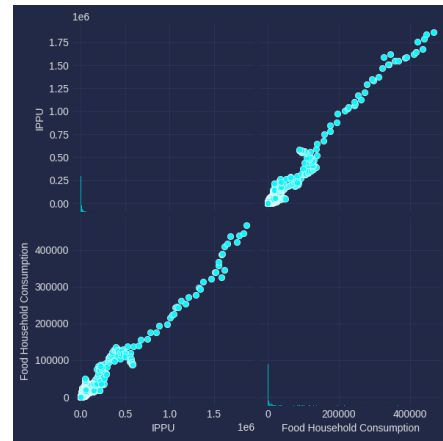


Fig. 3: *Pair Plot between features IPPU and Food Household Consumption*

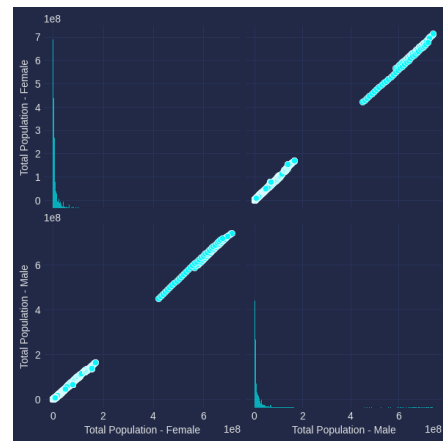


Fig. 4: *Pair Plot between features Total Population-Female and Total Population-Male*

and validation. These metrics provide insights into the accuracy and performance of the models.

D. Machine Learning Models Used

- Random Forest Regressor
- Linear Regressor
- Gradient Boosting Regressor
- Adaboost Regressor
- XGB Regressor
- Lasso Regressor
- Ridge Regressor
- Support Vector Regression (SVR) with various kernels: RBF kernel, Linear kernel, Polynomial kernel, Sigmoid kernel
- Artificial Neural Network (ANN): Two Architectures were designed, one for Label-encoded data (Fig. 5) and another for One-Hot encoded data (Fig. 6). The parameters that we used are the following:
 - 1) *optimizer*: Stochastic gradient descent
 - 2) *Number of Iterations*: 500
 - 3) *Learning rate*: $1e-4$ (Low because there were high variations in the output with higher lr.
 - 4) *Activation Function*: : ReLU

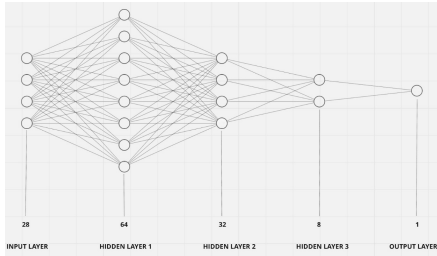


Fig. 5: ANN Arch. 1 (Label Encoding)

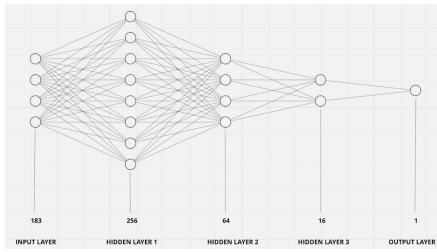


Fig. 6: ANN Arch. 2 (OneHot Encoding)

E. Improvement Methods

- **Outlier Removal**: removing the data points which deviate significantly from the rest of the data. These data points can deviate the model from learning therefore it's better to remove them. We have used IsolationForest for outlier detection.
- **Feature Engineering**: the creation of new features from the existing ones. The new feature added is the cluster-ID

feature containing the cluster number of the cluster the given data point belongs to. For determining the cluster IDs, we used the KMeans method with K equal to 3.

V. Result and Analysis

The results of various evaluation metrics for different models and preprocessing techniques (Label Encoding and One-Hot Encoding) are provided in Tables 1, 2, and 3. The evaluation metrics used include Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R2 Score for both training and validation sets.

A. Label-Encoded Data

Random Forest Regressor and XGBRegressor performed the best, showing low errors on both training and validation sets. Their high R2 score indicates they effectively captured the data patterns, highlighting their strong predictive ability. Support Vector Regression with Sigmoid Kernel is the worst performer, displaying extremely high training and validation errors and negative R2 scores on the training set (-17463.75), indicating an inappropriate fit. Moreover, the negative R2 score on the validation set (-18086.22) emphasizes the model's inability to learn from the data. The results are shown in Table 1.

TABLE I: Evaluation Metrics for Label-Encoded Data

Table Models	Training			Validation		
	MAE	RMSE	R2 Score	MAE	RMSE	R2 Score
RFR	0.103	0.141	0.940	0.276	0.375	0.581
LR	0.356	0.471	0.339	0.359	0.474	0.330
GBR	0.265	0.356	0.622	0.293	0.394	0.536
ABR	0.342	0.436	0.433	0.350	0.454	0.386
XGBR	0.088	0.116	0.959	0.277	0.373	0.585
LassoR	0.356	0.472	0.338	0.358	0.474	0.329
RidgeR	0.356	0.471	0.339	0.358	0.474	0.330
SVR (RBF)	0.330	0.446	0.408	0.342	0.456	0.379
SVR (Linear)	0.354	0.474	0.331	0.358	0.478	0.319
SVR (Poly)	0.357	0.482	0.309	0.403	1.139	-7.049
SVR (Sigmoid)	22.134	76.705	-17463.75	22.380	76.825	-18086.22

B. One-Hot Encoded Data

Random Forest Regressor and XGBRegressor continue to perform as the best models in our analysis. There isn't a significant difference observed between the results of models on label encoding and one-hot encoded data. One-hot encoding typically leads to a larger feature space due to the creation of binary columns for categorical data. However, this increase in feature complexity has not yielded better results for our models; instead, it has added complexity without a corresponding improvement in performance. With the increased number of features, Linear regression becomes the worst performer. The huge difference between the training and validation scores shows that it overfits the data. The results are shown in table 2.

TABLE II: Evaluation Metrics for One-Hot Encoded Data

Table	Training			Validation		
Models	MAE	RMSE	R2 Score	MAE	RMSE	R2 Score
RFR	0.102	0.141	0.940	0.277	0.375	0.580
LR	0.295	0.399	0.525	2.322e+8	4.952e+9	-1.77e+20
GBR	0.272	0.364	0.604	0.297	0.398	0.527
ABR	0.345	0.438	0.428	0.352	0.457	0.377
XGBR	0.109	0.144	0.937	0.277	0.374	0.583
LassoR	0.294	0.400	0.524	0.308	0.418	0.478
RidgeR	0.294	0.399	0.525	0.308	0.418	0.478
SVR (RBF)	0.284	0.397	0.531	0.307	0.421	0.472
SVR (Linear)	0.292	0.404	0.515	0.310	0.422	0.469
SVR (Poly)	0.316	0.441	0.422	0.344	0.469	0.343
SVR (Sigmoid)	1.217	5.682	-95.385	1.296	5.771	-107.049

C. One-Hot Encoded Data with PCA

On performing PCA on the one-hot encoded data, the performance of most of the models is reduced. The results are shown in Table 3. There is a significant improvement in the results of Linear Regression by reducing some features because the model becomes less prone to overfitting, allowing it to generalize better to unseen data. However, with the reduced performance of our best models, it is not preferred to implement PCA in this particular case, as it negatively impacts the overall predictive power of our models. The results are shown in Table 3.

TABLE III: Evaluation Metrics for One-Hot Encoded Data with PCA

Table	Training			Validation		
Models	MAE	RMSE	R2 Score	MAE	RMSE	R2 Score
RFR	0.116	0.159	0.924	0.312	0.425	0.461
LR	0.296	0.403	0.517	0.309	0.420	0.472
GBR	0.275	0.364	0.605	0.311	0.419	0.476
ABR	0.345	0.441	0.421	0.357	0.468	0.347
XGBR	0.118	0.156	0.927	0.326	0.444	0.413
LassoR	0.296	0.403	0.517	0.309	0.420	0.473
RidgeR	0.296	0.403	0.517	0.309	0.420	0.472
SVR (RBF)	0.285	0.398	0.529	0.307	0.422	0.471
SVR (Linear)	0.294	0.407	0.507	0.311	0.424	0.464
SVR (Poly)	0.318	0.442	0.417	0.345	0.473	0.334
SVR (Sigmoid)	1.239	5.752	-97.886	1.318	5.844	-110.239

D. Outlier Removal

The results of various evaluation metrics of the Random Forest and XGB regressors trained on the Label encoded and outlier removed dataset are shown in Table 4. Our finding reveals that outlier removal actually reduces the performance of the models, so it didn't really improve the performance.

E. Feature Engineering (Cluster ID)

In an attempt to enhance the performance of the Random Forest Regressor and XGBoost Regressor on label-encoded data,

TABLE IV: Evaluation Metrics for Label Encoded Data after Outlier Removal

Table	Training			Validation		
Models	MAE	RMSE	R2 Score	MAE	RMSE	R2 Score
RFR	0.104	0.142	0.941	0.283	0.385	0.560
XGBR	0.09	0.117	0.96	0.288	0.444	0.551

we introduced an additional feature into the dataset—Cluster ID—with the expectation that this inclusion would lead to improved performance. The results are detailed in Table 5. However, the outcomes indicate that this method also falls short in improving their performance. There is not much difference in the performance.

TABLE V: Evaluation Metrics for Label Encoded Data with Feature Engineering (Cluster ID)

Table	Training			Validation		
Models	MAE	RMSE	R2 Score	MAE	RMSE	R2 Score
RFR	0.103	0.14	0.941	0.277	0.375	0.582
XGBR	0.088	0.117	0.96	0.278	0.373	0.585

F. ANNs

We trained two ANN models, one is for label-encoded data and the other for one hot-encoded data. The results are shown in Table 6 and the Arch. of the ANNs are shown in Fig. 5 and Fig. 6. The results show that the performance of both the models is not good enough compared to other models. This suggests that overly complex models did not perform well in this context.

TABLE VI: Evaluation Metrics of ANNs

Table	Training			Validation		
Architecture	MAE	RMSE	R2 Score	MAE	RMSE	R2 Score
Arch. 1 (Label enc.)	0.313	0.416	0.483	0.334	0.444	0.427
Arch. 2 (OneHot enc.)	0.266	0.364	0.605	0.304	0.427	0.47

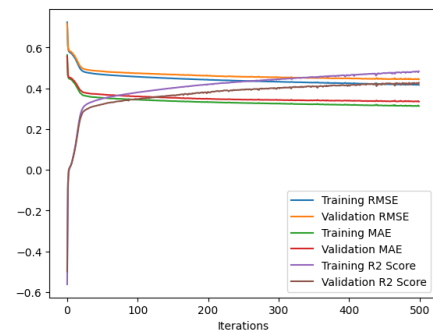


Fig. 7: Evaluation Metrics of Arch. 1 vs Iterations

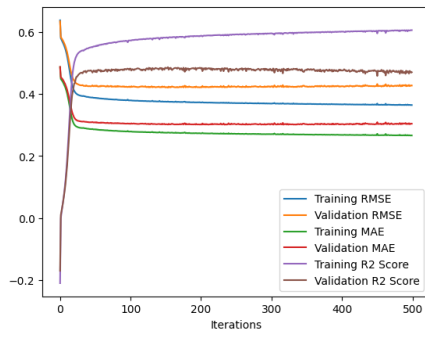


Fig. 8: *Evaluation Metrics of Arch. 2 vs Iterations*

VI. Conclusion

Our analysis of the data, training of various linear models, ensemble models, ANNs, etc, and various encoding methods led to the conclusion that the Random Forest regressor and XG Boost Regressor model with label encoding preprocessing on the categorical features performed well for our Agri-food dataset [1]. Although these models depicted a case of overfitting by showing high variance and low bias (as observed by the R2 scores from the training and validation sets), their R2 score on the validation set was still substantially better than those observed from other models on the validation set. Our analysis also showed that using either Label Encoding or One-Hot Encoding doesn't make a big difference in the results. It also explains the "Curse of dimensionality" by showing how a simple model like Linear regression's performance reduces drastically with an increase in the dimensions. We tried various Improvement methods although none of them really helped much. In essence, choosing the right model is crucial in Machine Learning and our focus was on exploring sophisticated models and fine-tuning them.

References

- [1] Agri-food CO2 emission dataset - Forecasting ML
- [2] Mahmoud Y. Shams et al. "A Machine Learning-Based Model for Predicting Temperature Under the Effects of Climate Change". DOI: 10.1007/978-3-031-22456-0_4
- [3] S. Salcedo-Sanz et al. "Monthly prediction of air temperature in Australia and New Zealand with machine learning algorithms". DOI: 10.1007/s00704-015-1480-4
- [4] Himanshu Vishwakarma. "Climate Change Analysis Using Machine Learning". DOI: 10.21275/SR20722101621